




1

Outline


- What is Reinforcement Learning (RL)
- Examples
- Defining an RL problem
 - Markov Decision Processes
- Solving an RL problem
 - Dynamic Programming
 - Monte Carlo methods
 - Temporal-Difference learning




2




3




4




5




6




7




8




9



10



11



12

Is RL really necessary?

Would supervised learning do the job?

13

Is RL really necessary?

How about we turn it into a supervised learning problem?

14

Is RL really necessary?

How about we turn it into a supervised learning problem?

Two major issues:

1. How to create these datasets for each task? (e.g., Pong, Breakout, Space Invaders, etc.)
2. Who will provide the labels for these Supervised Learning (SL) datasets? (e.g., best gamers??)

2a. Will the machine ever beat humans if its actions can only (at best) match a human player (who provides the labels)?

15

Characteristics of Reinforcement Learning

What makes reinforcement learning different from other machine learning paradigms?

- There is no supervisor, only a reward signal
- Feedback is delayed, not instantaneous
- Time really matters (sequential, non i.i.d data)
- Agent's actions affect the subsequent data it receives

Source: David Silver (<https://www.davidsilver.uk/teaching/>)

16

Outline

- Examples
- Defining an RL problem
 - Markov Decision Processes
- Solving an RL problem
 - Dynamic Programming
 - Monte Carlo methods
 - Temporal-Difference learning

17

What is Reinforcement Learning

Learning (by operating in and interacting with the world [aka. environment]) to make good sequence of decisions under uncertainty

- Reward/Penalty
 - Reward function
 - Value function
- Sequential decision making
 - Policy
- Environment
 - States
 - Uncertainty
- Learning
 - Observation
 - Algorithms
 - Exploration vs Exploitation

18

What is Reinforcement Learning - Reward

- **Reward** (at time t) R_t is a numerical feedback signal
- Indicates how well agent is doing at step t
- The agent's job is to maximise **cumulative reward**
 - Sequence of decisions → sequence of actions
- Rewards can be **delayed**
 - Delayed rewards can be “**discounted**”
 - See next slide on “**Sequential decision making**”


19

What is Reinforcement Learning – Sequential decision making

- **Goal:** select actions to maximise total future reward
- Actions may have long term (i.e., delayed) consequences
- It may be better to sacrifice immediate reward to gain more long-term reward
 - Study (vs party) now
 - Contribute to your superannuation (less money to spend now vs retire comfortably)
 - Self-driving EV stops now to recharge (vs stranded in the middle of nowhere)

20

Example - Robot in a room



actions: UP, DOWN, LEFT, RIGHT

Deterministic vs
Stochastic

- **reward** +100 at [4,3], -100 at [4,2] (and they are also terminal states)

21

Example - Robot in a room




actions: UP, DOWN, LEFT, RIGHT

UP
80%
10%
10%move UP
move LEFT
move RIGHT

- **reward** +100 at [4,3], -100 at [4,2] (and they are also terminal states)

22

Example - Robot in a room




actions: UP, DOWN, LEFT, RIGHT

UP
80%
10%
10%move UP
move LEFT
move RIGHT

- **reward** +100 at [4,3], -100 at [4,2] (and they are also terminal states)
- **reward** -4 for each step

23



24

Agent and Environment

Source: David Silver (<https://www.davidsilver.uk/teaching/>)

- At each step t the agent:
 - Executes action A_t
 - Receives observation O_t
 - Receives scalar reward R_t
- The environment:
 - Receives action A_t
 - Emits observation O_{t+1}
 - Emits scalar reward R_{t+1}
- t increments at env. step

25

History and State

Source: David Silver (<https://www.davidsilver.uk/teaching/>)

- The **history** is the sequence of observations, actions, rewards

$$H_t = O_0, R_0, A_0, \dots, A_{t-1}, O_t, R_t$$
- i.e. all observable variables up to the current time t
- What happens next depends on the history:
 - The agent selects actions
 - The environment selects observations/rewards
- State** is the information used to determine what happens next
 - i.e., state is a function of the history:

$$S_t = f(H_t)$$

26

Environment State

Source: David Silver (<https://www.davidsilver.uk/teaching/>)

- The **environment state** S^e_t is the environment's private representation
- The environment state is not usually visible to the agent
- Even if S^e_t is visible, it may contain irrelevant information

27

Agent State

Source: David Silver (<https://www.davidsilver.uk/teaching/>)

- The **agent state** S^a_t is the agent's internal representation (of the environment)
- i.e. whatever information the agent uses to pick the next action
- It can be any function of the history:

$$S^a_t = f(H_t)$$

28

Information State

Source: David Silver (<https://www.davidsilver.uk/teaching/>)

- An **information state** (a.k.a. **Markov state**) contains all useful information from the history.

Definition: A state S_t is Markov if and only if
$$P[S_{t+1} | S_t] = P[S_{t+1} | S_1, \dots, S_t]$$

- i.e., "*The future is independent of the past given the present*"
- Existing RL algorithms assumes that the agent state S^a_t is Markov.

29

Fully Observable Environment

Source: David Silver (<https://www.davidsilver.uk/teaching/>)

- Full observability:** agent directly observes environment state

$$O_t = S^a_t = S^e_t$$
- Agent state = environment state = information state
- Formally, this is a **Markov Decision Process (MDP)**

30


Partially Observable Environment

Source: David Silver
(<https://www.davidsilver.uk/teaching/>)

- Partial observability: agent indirectly observes environment with its sensors:
 - Its sensors can be incomplete
 - Its sensors can be noisy
 - There are other agents who have private information not available to the RL agent (e.g., agent playing card games, warfare, etc.)
- Now agent state \neq environment state
- Formally this is a partially observable Markov decision process (POMDP)
- Agent must construct its own state representation S^a_t :
 - e.g. Complete (observable) history: $S^a_t = H_t$

31

RL Agents – Making decisions and Learning




Major Components of an RL Agent:

- **Policy:** (Mandatory) agent's behaviour function
- **Value function:** (Optional) how good is each state and/or action
- **Model:** (Optional) agent's representation of the environment

32


Robot in a room (or, Mouse in a maze)



- reward +100 at [4,3] (the cheese), -100 at [4,2] (the trap)
- reward -4 for each step
- For an MDP, the states consist of the whole maze (i.e., agent having a bird's-eye view of the maze)
- what's the strategy to achieve max reward?

33


Is this a solution?



- only if actions are deterministic
 - not in this case (actions are stochastic)
- Solution is a **policy**
 - **Policy:** mapping from each state to an action


34

Which policy?



- Deterministic policy: $a = \pi(s)$
- Stochastic policy: $\pi(a | s) = P[A_t = a | S_t = s]$


Note: Deterministic policy \neq
Deterministic environment



35

RL algorithms - Value function

- Assume for a moment that the environment is DETERMINISTIC, optimal policy:



- We can calculate (recursively from the terminal states) the following value function:

88	92	96	+100
84		92	-100
80	84	88	84

36

RL algorithms - Value function

- Value function is a prediction of future reward
- Used to evaluate the goodness/badness of states
- And therefore to select between actions
- Example:

$$v_{\pi}(s) = \mathbb{E}_{\pi} [R_{t+1} + \gamma R_{t+2} + \gamma^2 R_{t+3} + \dots | S_t = s]$$

37



What about STOCHASTIC environments!

Optimal policy: - WHY?

			+100
			-100

UP

- 80% move UP
10% move LEFT
10% move RIGHT

Can you calculate the **value function**?

HINT: It is the **expected value** based on the probabilistic outcomes of the action taken in each cell!

38



Reward for each step: -120



			+100
			-100

39

Reward for each step: -10



			+100
			-100

40

Reward for each step: -4



			+100
			-100

41

Reward for each step: -1



			+100
			-100

42

Reward for each step: +1

43

Markov Decision Process (MDP) - Model

- set of **states** S , set of **actions** A , **initial state** s_0
- **transition** model $P(s'|s,a)$
 - $P([1,2] | [1,1], \text{up}) = 0.8$
- **reward function** $r(s)$
 - $r([4,3]) = +100$

44

Categorizing RL agents (1)

Source: David Silver (<https://www.davidsilver.uk/teaching/>)

- Value Based
 - No Policy (Implicit)
 - Value Function
- Policy Based
 - Policy
 - No Value Function
- Actor Critic
 - Policy
 - Value Function

45

Categorizing RL agents (2)

Source: David Silver (<https://www.davidsilver.uk/teaching/>)

- Model Free
 - Policy and/or Value Function
 - No Model
- Model Based
 - Policy and/or Value Function
 - Model

46

RL Agent Taxonomy

Source: David Silver (<https://www.davidsilver.uk/teaching/>)

47

Computing return from rewards

- episodic (vs. continuing) tasks
 - "game over" after N steps
 - optimal policy depends on N ; harder to analyze
- additive rewards
 - $V(s_0, s_1, \dots) = r(s_0) + r(s_1) + r(s_2) + \dots$
 - infinite value for continuing tasks
- discounted rewards ($\gamma < 1$)
 - $V(s_0, s_1, \dots) = r(s_0) + \gamma * r(s_1) + \gamma^2 * r(s_2) + \dots$
 - value bounded if rewards bounded

48

Value functions

- state value function: $V^\pi(s)$
 - expected return when starting in s and following π
- state-action value function: $Q^\pi(s, a)$
 - expected return when starting in s , performing a , and following π
- For a fixed policy π , the utility function obey the **Bellman equation**:

$$V^\pi(s) = r(s) + \gamma \sum_{s'} P(s' | s, \pi(s)) * V^\pi(s')$$

49

Value functions

- state value function: $V^\pi(s)$
 - expected return when starting in s and following π
- state-action value function: $Q^\pi(s, a)$
 - expected return when starting in s , performing a , and following π
- For a fixed policy π , the utility function obey the **Bellman equation**:

$$V^\pi(s) = r(s) + \gamma \sum_{s'} P(s' | s, \pi(s)) * V^\pi(s')$$

50

Optimal value functions

- there's a set of *optimal* policies
 - V^* defines the value of policy for every state $s \in S$
 - Then, there is a single **optimal value** for each state $s \in S$
- Bellman optimality equation**
 - $V^*(s) = r(s) + \gamma * \max_{a \in A(s)} \sum_{s'} P(s' | s, a) * V^*(s')$
 - system of n non-linear equations
 - solve for $V^*(s)$
 - easy to extract the **optimal policy** π^*
- having $Q^*(s, a)$ makes it even simpler: $\pi^*(s) = \arg \max_a Q^*(s, a)$

51

Outline

- Examples
- Defining an RL problem
 - Markov Decision Processes
- Solving an RL problem
 - Dynamic Programming
 - Monte Carlo methods
 - Temporal-Difference learning

52

Dynamic programming (DP)

- main idea
 - If an action in a situation/state cause something (really) good/bad to happen immediately then we "know" the "goodness" of that action (given that state) and the "value" of that state!
 - need a **perfect** model of the environment (aka. **Model-based RL**)
- two main components (of "**Policy iteration algorithm**")
 - policy evaluation: compute V^π from π
 - policy improvement: improve π based on V^π
- start with an arbitrary policy
- repeat evaluation/improvement until convergence

			+100
			-100
START			

53

Policy iteration

- Initialise V and policy π' for all $s \in S$

$\pi = \pi'$

For all $s \in S$: (concurrent update \odot)

$$V(s) = r(s) + \gamma \sum_{s'} P(s' | s, \pi(s)) * V(s')$$

For all $s \in S$:

$$\pi'(s) = \arg \max_{a \in A(s)} \sum_{s'} P(s' | s, a) * V(s')$$

Repeat until $\pi = \pi'$

54

Maintaining exploration

- deterministic/greedy policy won't explore all actions
 - don't know anything about the environment at the beginning
 - need to try all actions to find the optimal one
- maintain exploration
 - use soft policies instead: $\pi(s,a) > 0$ (for all s,a)
 - i.e., the policy does not simply return an action for each state, it returns a probability distribution over the action space for each state
- ϵ -greedy policy
 - with probability $1-\epsilon$ perform the optimal/greedy action
 - with probability ϵ perform a random action
 - will keep exploring the environment
 - slowly move it towards greedy policy: $\epsilon \rightarrow 0$

61

Summary of Monte Carlo

- don't need model of environment
 - averaging of sample returns
 - only for episodic tasks
- learn from sample episodes or simulated experience
- can concentrate on "important" states
 - don't need a full sweep
- need to maintain exploration
 - use soft policies

62

Outline

- examples
- defining an RL problem
 - Markov Decision Processes
- solving an RL problem
 - Dynamic Programming
 - Monte Carlo methods
 - Temporal-Difference learning

63

Temporal Difference Learning

- combines ideas from MC and DP
 - like MC: learn directly from experience (don't need a model)
 - like DP: learn from values of successors
 - works for continuous tasks, usually faster than MC
- constant-alpha MC:
 - have to wait until the end of episode to update
$$V(s_t) \leftarrow V(s_t) + \alpha [R_t - V(s_t)]$$
- simplest TD
 - update after every step, based on the successor
$$V(s_t) \leftarrow V(s_t) + \alpha [r_{t+1} + \gamma V(s_{t+1}) - V(s_t)]$$

64

Sarsa

- again, need $Q(s,a)$, not just $V(s)$

$$Q(s_t, a_t) \leftarrow Q(s_t, a_t) + \alpha [r_t + \gamma Q(s_{t+1}, a_{t+1}) - Q(s_t, a_t)]$$

- control
 - start with a random policy
 - update Q and π after each step
 - again, need ϵ -soft policies

65

The RL Intro book

Richard Sutton, Andrew Barto
Reinforcement Learning,
An Introduction

<http://www.cs.ualberta.ca/~sutton/book/the-book.html>

66

Summary

- Reinforcement learning
 - use when need to make decisions in uncertain environment
- solution methods
 - dynamic programming
 - need complete model
 - Monte Carlo
 - time-difference learning (Sarsa, Q-learning)
- most work
 - algorithms simple
 - need to design features, state representation, rewards



Where to start?

- You need an environment (to build your RL agent)
- The OpenAI Gym provides a good starting point with some pre-defined environments and also allow you to define your own environment:
- <https://towardsdatascience.com/reinforcement-learning-with-openai-d445c2c687d2>

67

68